FigJam Board Generation Prompt for Claude.ai

Instructions: Copy everything below and paste into Claude.ai with Figma Connector enabled

---

PROMPT START

Create a comprehensive, well-organized FigJam board titled "Craftit AI — Research & Strategy Board 🐙" using ALL of the following research data. This is a product strategy board for Craftit AI, an AI-first image and video generation tool.

Organize the board into clearly separated sections with headers, using sticky notes, shapes, connectors, and visual hierarchy. Use color coding consistently throughout:

- 🔴 Red = Pain points, risks, threats
- 🟢 Green = Opportunities, strengths
- 🟡 Yellow = Insights, findings
- 🔵 Blue = Strategy, actions
- 🟣 Purple = Personas, users
- 🟠 Orange = Competitors

SECTION 1: THE BIG INSIGHT (Center top, large and prominent)

Create a large, prominent section at the top center:

"The market is screaming for an AI-first Canva — a tool that combines Canva's simplicity with Midjourney's quality and Runway's video capability. Nobody has built it yet. Craftit AI has the feature set to be that product."

SECTION 2: MARKET STATS (6 key numbers in a row)

Create 6 stat cards side by side:

1. **$50-80B** — AI Creative Tools TAM by 2028
2. **45-55% CAGR** — AI Video Generation growth (fastest segment)
3. **260M MAU** — Canva proves "design for everyone" works at scale
4. **$500M/yr** — Midjourney revenue from image gen alone
5. **~0** — AI creation tools built for kids (massive gap)
6. **3-5 tools** — Average tools per user project (users hate this)

Market Trends (4 sticky notes below stats):

🚀 **AI Video at Inflection Point** Quality now good enough for social content but tools still complex. Simplified tool wins NOW. Video is 18-24 months behind image gen in quality/accessibility.

📱 **Mobile is Blue Ocean** SEA, India, LatAm = mobile-only markets. 3B+ smartphone users. Runway & Midjourney have NO mobile apps. Canva has best mobile but weak AI.

👶 **Kids = Zero Competition** ~50-100M kids interested in digital creation. ZERO purpose-built AI tools for under-13. COPPA compliance creates barrier that deters competitors. Canva Education has 60M+ students.

🔄 **Consolidation Coming** Hundreds of AI tools launched 2023-2025. "Middle squeeze" starting: mega-platforms + niche specialists survive. Generic tools die. Time to differentiate.

Market Sizing Detail:

- Overall Generative AI Market 2025: $37.89B globally (Precedence Research)
- 2035 projection: $1,206B at 36.97% CAGR
- Media & entertainment: 34%+ of generative AI revenue
- North America: 41% market share
- AI Image Generation estimated 2025: $5-8B, projected $15-20B by 2028 at ~30% CAGR
- AI Video Generation estimated 2025: $1.5-3B, projected $8-15B by 2028 at 45-55% CAGR
- Web Content Management market: $24.97B by 2029 at 18.6% CAGR

Revenue Models in the Space:

Model	Examples	Typical Pricing
Freemium + Subscription	Canva, Midjourney	$10-30/mo consumer, $50-150/mo pro
Credit/Token-based	DALL-E, Runway, Leonardo	Pay-per-generation
Enterprise/Team	Canva Teams, Adobe	$100-500+/seat/mo
API/Platform	Stability AI, OpenAI	Usage-based
Hybrid	Most leaders	Sub + credits (emerging best practice)

Enterprise GenAI budgets averaged $7M in 2023, plans to increase 2-5x in 2024. Spend shifting from "innovation budgets" to permanent software line items.

User Segments:

Segment	Est. Size (Global)	Growth	WTP
Professional Creators	15-25M	15-20%/yr	$20-50/mo
Marketers & Social Media	30-50M	25-30%/yr	$15-30/mo
SMBs & Businesses	40-60M teams	30-35%/yr	$30-100+/mo team
Casual Users / Hobbyists	200-400M	40-50%/yr	$0-10/mo
Students	60-100M	35-40%/yr	$0-5/mo
Kids (under 13)	50-100M potential	Emerging	$5-15/mo parent-funded

Conversion benchmarks: Free-to-paid 3-10% industry average. Canva: ~8% (21M paid / 260M MAU). Average consumer LTV: $120-360/year.

Geographic Analysis:

Region	Share	Key Characteristics
North America	~41%	Highest ARPU, enterprise-heavy, early adopter
Europe	~22%	GDPR/AI Act aware, quality-conscious, multilingual

Region	Share	Key Characteristics
Asia-Pacific	~28%	Mobile-first, price-sensitive, fastest growth 27.6% CAGR
Latin America	~5%	Emerging, mobile-dominant, creative culture
Middle East & Africa	~4%	Young population, mobile-only

Regional preferences: US/EU = quality + brand safety. China = domestic ecosystem. India = extremely price-sensitive, mobile-only. Japan/Korea = high design standards, anime. SEA = mobile-first, social commerce. LatAm = creator economy booming.

SECTION 3: COMPETITOR ANALYSIS (Large matrix section)

Competitor Feature Matrix (create as a table/grid):

Competitor	Image Gen	Video Gen	Canvas/Design	Long Video	Community	Mobile App	A S
Midjourney	★★★★★	×	×	×	★★★★	×	×
Runway	★★★	★★★★★	×	×	★★	×	×
Pika	×	★★★★	×	×	★★	★★★	×
Canva AI	★★	★★	★★★★★	×	×	★★★★★	×
Adobe Firefly	★★★	★★	★★★★ (via CC)	×	★★	★★★	×
Higgsfield AI	★★★	★★★★	×	★★★	★★★★★	★★	★★
Freepik	★★★	★★	★★★	×	×	★★	×
OpenArt AI	★★★★	★★★	×	×	★★★	★★	★★
Luma AI	★★★	★★★★★	×	×	★	★★★★	×
Kling AI	★★★	★★★★	×	×	★★	★★★	×
PixVerse	×	★★★★	×	★★	★★	★★	★★
Leonardo AI	★★★★	★★	★★★	×	★★★	★★★	×
FLUX	★★★★★	×	×	×	★★ (open source)	×	×

Competitor	Image Gen	Video Gen	Canvas/Design	Long Video	Community	Mobile App	A S
Ideogram	★★★★	×	★★	×	★★	×	×
Haiper AI	×	★★★	×	×	★	×	×
Vidu	×	★★★★	×	★★	★	★★	×
CRAFTIT AI	★★★★	★★★★	★★★★ ➡ SOON	★★★★ ➡ SOON	★★★	?	★★ ➡ SOON

Individual Competitor Deep Dives (create as grouped sticky notes per competitor):

1. Midjourney

- Core: Best-in-class image generation (V6+), web editor with inpainting/variation/zoom. No native video. No canvas/design tools. Recently added web app (beyond Discord-only).
- Pricing: Basic $10/mo (200 gens), Standard $30/mo, Pro $60/mo, Mega $120/mo. ~20% annual discount.
- Target: Creators, artists, designers. Leans pro-creative but attracts casuals.
- UX: Originally Discord-only (high friction). Web app now available but still prompt-heavy. No guided workflows.
- Strengths: Industry-leading aesthetics, massive brand recognition, huge Discord community (~19M members).
- Weaknesses: No video at all. No design/canvas tools. No mobile app. Discord UX alienates non-technical users. No templates. Limited editing.
- Community: Huge Discord (~19M). Showcase channels. No native in-app feed.
- Recent: Launched web editor, exploring hardware. Added editing features. Still image-only.
- Mobile: Web-based only, no native app.
- **Relevance to Craftit:** No video, no canvas. Craftit differentiates on full image→video→design pipeline and simplicity.

2. Runway

- Core: Gen-4/Gen-4.5 video (text/image-to-video), Gen-4 image gen, video editing suite, Act-Two (performance capture), Veo integration, TTS, audio gen, upscaling. Most complete pro video platform.
- Pricing: Free (125 one-time credits), Standard $12/mo (625 credits), Pro $28/mo (2,250), Unlimited $76/mo.
- Target: Professional filmmakers, video producers, creative agencies. Increasingly prosumers.
- UX: Powerful but complex. Multiple tools, credit system, model selection, parameters. Not beginner-friendly.
- Strengths: Best-in-class video quality. Full editing suite. Enterprise features. API. Multi-model access. Performance capture unique.
- Weaknesses: Expensive for casuals. Complex UX. Credits burn fast. No design/canvas. No templates for non-video. Steep learning curve.
- Community: Discord. No in-app social feed.
- Recent: Gen-4.5, Aleph video editing, Act-Two, integrating Veo. Pushing toward full filmmaking platform.
- Mobile: No dedicated app. Web only.
- **Relevance to Craftit:** Pro video leader but leaves casuals behind. Craftit's simplicity-first + auto-mode video directly targets Runway's complexity gap.

3. Pika

- Core: Video gen (text/image-to-video), Pika Effects (lip sync, sound effects, modify region), scene extension. Focused on short video.
- Pricing: Free tier. Standard ~$8/mo, Pro ~$28/mo, Unlimited ~$58/mo (credit-based).
- Target: Social media creators, casual users, younger demographic.
- UX: Simple and playful. Clean interface. Low barrier.

- Strengths: Very accessible UX. Fun viral effects (lip sync). Good free tier. Mobile-friendly. Strong social media presence.
- Weaknesses: Video quality below Runway/Luma for cinematic. No standalone image gen. No design tools. No canvas. Limited editing. No long video.
- Community: Discord. Social sharing around viral TikTok/X effects.
- Recent: Pika 2.0/2.1 improvements. Doubling down on viral effects and accessibility.
- Mobile: Mobile-optimized web. Likely iOS app.
- **Relevance to Craftit:** Closest competitor in "simple video gen" space. Craftit needs to match Pika's simplicity while offering more (image gen, canvas, long video).

4. Canva AI

- Core: Full design suite (posters, presentations, social, videos), AI image gen (Magic Media), AI video gen, text effects, background removal, Magic Eraser/Expand, millions of templates, brand kits, collaboration.
- Pricing: Free (limited AI), Pro $13/mo/person (500 AI credits), Teams $10/mo/person, Enterprise custom.
- Target: Everyone — businesses, marketers, students, casual. "Design for everyone."
- UX: Extremely simple. Template-first. Drag-and-drop. Gold standard for accessible design tools.
- Strengths: Massive templates. Full design ecosystem. Brand kits. Collaboration. Print. 170M+ MAU. AI integrated into existing workflows. Trust.
- Weaknesses: AI generation quality mediocre vs dedicated tools. Video AI basic. Not AI-first — AI bolted on. Image gen not competitive with MJ/DALL-E. No community showcase.
- Community: No creator community. Template marketplace from designers.
- Recent: Heavy AI integration. Magic Studio. Acquiring AI capabilities. Pushing enterprise.
- Mobile: Excellent app (iOS/Android). Full-featured.
- **Relevance to Craftit:** BIGGEST STRATEGIC OVERLAP AND OPPORTUNITY. Canva = canvas king with weak AI. Craftit's upcoming canvas + superior AI gen could capture users who want Canva's ease but AI-quality output.

5. Adobe Firefly

- Core: Image gen (text-to-image), generative fill/expand, text effects, style transfer, vector gen. Integrated into Photoshop/Illustrator/Premiere. Video gen (Firefly Video) recently added.
- Pricing: Free (25 credits/mo), Premium $9.99/mo (2,000 credits), included in Creative Cloud.
- Target: Professional designers, photographers, video editors in Adobe ecosystem. Enterprise.
- UX: Web app clean and simple. Full power requires PS/AI knowledge. Dual-track: simple web + powerful desktop.
- Strengths: Commercially safe (licensed training data). Deep CC integration. Enterprise trust. Content Credentials. Professional editing + AI.
- Weaknesses: Image quality lags MJ/FLUX. Video gen new and behind Runway/Pika. Conservative approach. Expensive with full CC. Not standalone enough.
- Community: Behance (separate). No in-app community for Firefly.
- Recent: Firefly Video model, deeper CC integration, Firefly Services API, content credentials.
- Mobile: Adobe Express mobile includes Firefly. PS mobile has some AI.
- **Relevance to Craftit:** Targets pros in Adobe ecosystem. Craftit wins users who want AI-quality output without PS skills or CC subscription.

6. Higgsfield AI

- Core: AI video gen, image gen (Soul 2.0), Cinema Studio 2.0 (3D scene control, cinematic sequences), Vibe Motion (prompt→motion design), AI Influencer studio, character consistency, contests, integrated Kling 3.0, chat assistant.
- Pricing: Freemium with credit tiers. Runs contests with $500K prizes.
- Target: Content creators, social media influencers, AI art enthusiasts. Younger, trend-forward.
- UX: Feature-rich but approachable. Multiple specialized studios. Community-driven with contests.
- Strengths: Unique Cinema Studio with 3D scene control. Multi-model integration (Kling). AI Influencer studio differentiated. Large contest prizes drive engagement. Active community.

- Weaknesses: Brand awareness much lower than Runway/MJ. Feature sprawl — many studios may confuse. Quality depends on third-party models. No design/canvas.
- Community: Built-in community with contests, trending, sharing. Strongest community features among video-gen competitors.
- Recent: Cinema Studio 2.0, Soul 2.0, Kling 3.0 integration, $500K contests. Aggressive growth through community.
- Mobile: Web-first. May have mobile optimization.
- **Relevance to Craftit:** Community/contest model worth studying. Cinema Studio similar to Craftit's manual timeline. Auto-mode long video could leapfrog Cinema Studio's manual approach.

7. Freepik

- Core: Massive stock library (photos, vectors, icons, videos, PSD), AI image gen (Mystic, FLUX integration), AI video gen, mockup gen, design editor, presentation maker.
- Pricing: Free (limited), Essential ~$9/mo, Plus ~$15/mo, Premium ~$25/mo.
- Target: Designers, marketers, small businesses needing stock + quick design.
- UX: Stock-library-first with AI as enhancement. Simple search→download. AI tools secondary.
- Strengths: Enormous asset library. Established brand. AI as complement to stock. Affordable. Good for everyday business design.
- Weaknesses: AI gen not best-in-class. Not AI-first. Video gen basic. No AI community. Design editor limited vs Canva.
- Community: Contributor marketplace. No social/sharing for AI.
- Recent: AI across all asset types. Mystic model. FLUX integration. Moving from stock → AI-augmented.
- Mobile: Apps available but download-focused.
- **Relevance to Craftit:** Competes on stock+AI hybrid. Craftit is AI-first = fundamentally different. Low direct threat but affordable pricing + stock library compete for budget-conscious users.

8. OpenArt AI

- Core: AI image gen (100+ models including FLUX, SDXL, MJ-style), AI video gen, consistent character creation, personalized model training, One-Click Story, audio models, workflows. Multi-model platform.
- Pricing: Free (40 trial credits), Essential $7/mo (4,000 credits), Advanced $14.50/mo (12,000), Infinite $28/mo (24,000), Wonder $120/mo. Team $17.50/seat. 50% annual discounts.
- Target: AI art enthusiasts, creators wanting model variety, hobbyists to semi-pros.
- UX: Model-picker approach — choose from 100+ models. Powerful but overwhelming. One-Click Story simplifies narrative.
- Strengths: Massive model selection. Competitive pricing. Character consistency. Story creation. Good credit value. Team plans.
- Weaknesses: Jack-of-all-trades, master of none. Quality depends on model choice (inconsistent). Cluttered UI. Less known brand. No design/canvas.
- Community: Has community features but less prominent than MJ Discord.
- Recent: One-Click Story, consistent characters, expanding models, audio gen. Moving toward narrative content.
- Mobile: Web-based, mobile-responsive.
- **Relevance to Craftit:** One-Click Story similar to planned auto-mode. Offers more models but less polish. Craftit should ensure auto-mode story experience is dramatically simpler and more integrated.

9. Luma AI (Dream Machine)

- Core: Video gen (Ray3/Ray3.14 — native 1080p, HDR), image gen (Photon), video modification, extend, references, upscaling, reframing. API available.
- Pricing: Free (limited, draft quality), Lite $7.99/mo (3,200 credits), Plus $23.99/mo (10,000), Unlimited $75.99/mo, Enterprise.
- Target: Pro video creators, filmmakers, API developers.
- UX: Clean, focused. Model-centric. Technical users who understand resolution/HDR.
- Strengths: Ray3 top-tier video quality. Native 1080p + HDR. Good API. Reasoning capabilities. Research-driven.
- Weaknesses: Expensive at high quality (HDR burns 4-8x credits). No design tools. No community. Complex credits.

- Community: Minimal. Discord exists but not central.
- Recent: Ray3.14 (4x faster, 3x cheaper, native 1080p). Pushing cinematic quality. HDR unique.
- Mobile: iOS app available.
- **Relevance to Craftit:** Pro video competitor but different audience. Craftit targets everyone; Luma targets filmmakers.

10. Kling AI (Kuaishou)

- Core: Video gen (Kling 3.0), image gen, lip sync, motion brush. Chinese-developed, global expansion.
- Pricing: Free tier generous. Pro from ~$8/mo. Aggressive pricing.
- Target: Global creators, especially strong in Asian markets.
- UX: Straightforward. Clean.
- Strengths: Very competitive video quality (3.0). Aggressive pricing. Strong R&D from Kuaishou. Good motion.
- Weaknesses: Brand trust outside Asia. Data privacy concerns. No design tools. Limited ecosystem.
- Community: Growing, especially via platforms like Higgsfield.
- Recent: Kling 3.0 public release. Expanding global access.
- Mobile: Apps available.

11. PixVerse

- Core: Video gen (V5.6), native audio gen (SFX, music, dialogue), multi-shot storytelling, text/image-to-video, real-time interactive world engine (research).
- Pricing: Freemium, credit-based. Affordable.
- Target: Creators wanting quick, complete videos with audio.
- UX: Simple. Full-screen creation. One-click storytelling.
- Strengths: Native audio gen unique — videos come with sound. Multi-shot storytelling built in. Fast 1080p. Strong cost-performance.
- Weaknesses: Less known. No design tools. Research-stage features.
- Community: Growing but small.
- Recent: V5.6 with audio-visual consistency. Real-time world engine research. Auto multi-shot.

12. FLUX (Black Forest Labs)

- Core: Image gen only (FLUX.1 Pro, Dev, Schnell). Text rendering excellence. Open-source models.
- Pricing: Starter $6.90/mo (500 credits), Premium $9.90/mo (2,000), Pro $27.90/mo (8,000). Also via third-party platforms.
- Target: Developers, AI enthusiasts, users wanting MJ-quality with open flexibility.
- UX: Simple web app. Also API and open-source.
- Strengths: Excellent quality. Best text rendering. Open-source. Available everywhere. Affordable.
- Weaknesses: Image only. No video. No design. Minimal own platform.
- Community: Open-source community. No social.
- Recent: FLUX 1.1 Pro improvements. Focus on quality/speed.

13. Leonardo AI

- Core: Image gen (Phoenix), real-time canvas, motion gen (image-to-video), texture gen for 3D, fine-tuned models, image editing.
- Pricing: Free (150 tokens/day), Apprentice ~$10/mo, Artisan ~$24/mo, Maestro ~$48/mo.
- Target: Game devs, digital artists, content creators.
- UX: Feature-rich with real-time canvas. More technical than Canva, more visual than ComfyUI.
- Strengths: Real-time canvas powerful. Good model training. 3D texture unique. Solid community. Good free tier.
- Weaknesses: Spread across too many features. Video basic. Not design-focused. Overwhelming.
- Community: Active, shared creations.
- Recent: Phoenix model, canvas upgrades, video features.

- Mobile: iOS app.

14. Ideogram

- Core: Image gen with industry-best text rendering, canvas/editing, style mixing.
- Pricing: Free tier, Plus ~$8/mo, Pro ~$20/mo.
- Target: Designers needing text in images, marketers, social media.
- UX: Clean, simple. Good at text-in-images.
- Strengths: Best text-in-image rendering. Clean UX. Design-oriented.
- Weaknesses: Narrow focus. No video. Limited editing.
- Community: Growing gallery.
- Recent: Ideogram 3.0 with improved realism and canvas.

15. Haiper AI

- Core: Video gen, animation. Founded by ex-DeepMind researchers.
- Pricing: Freemium. Affordable.
- Target: Casual creators wanting simple video gen.
- UX: Very simple. Minimal.
- Strengths: Simplicity. Free. Research pedigree.
- Weaknesses: Quality behind leaders. Limited features. Small team. No design.
- Community: Small.
- Recent: Haiper 2.0 improvements.

16. Vidu (Shengshu AI)

- Core: Video gen (Chinese-developed), long video support, multi-subject consistency.
- Pricing: Freemium, affordable.
- Target: Chinese market, expanding globally.
- UX: Clean, straightforward.
- Strengths: Strong video quality. Long video support. Competitive with Kling.
- Weaknesses: Limited global presence. Language barriers. No design.
- Community: China-focused.
- Recent: Global expansion, quality improvements.

Simplicity vs Power Matrix (create as 2x2 diagram):

HIGH POWER quadrant: Runway, Adobe Firefly (with CC), Luma AI MID POWER quadrant: Higgsfield, Kling AI, Leonardo AI SIMPLE + MID POWER: Pika, Haiper, PixVerse SIMPLE + HIGH POWER (design): Canva AI COMPLEX + MID: OpenArt AI, FLUX, Midjourney (prompt-heavy)

Audience Segmentation (create as spectrum):

- Casual: Pika, Haiper, PixVerse
- Prosumer: Midjourney, OpenArt, Higgsfield, Leonardo, Ideogram
- Pro: Runway, Luma AI, Adobe Firefly
- Business: Canva AI, Adobe Firefly, Freepik
- Enterprise: Runway, Adobe, Canva

Watch Closest (red sticky notes):

1. **Canva AI** — if they improve AI gen quality, biggest threat (have UX + 260M users + canvas already)
2. **Pika** — closest in simplicity + video gen positioning

3. **PixVerse** — auto storytelling with native audio is innovative
4. **Higgsfield** — community model and Cinema Studio overlap

Where Nobody Competes (green sticky notes):

1. Full pipeline: image → canvas → video in ONE tool
2. Auto long video with editorial control at every step
3. Quality AI generation + design canvas combined
4. Mobile-first AI creation (quality + simple)
5. AI creation for kids

SECTION 4: GAP ANALYSIS (6 opportunity cards)

Gap 1: The "Full Pipeline" Gap — BIGGEST OPPORTUNITY No competitor offers AI image gen → design canvas → short video → long video in one platform.

- Midjourney: image only
- Runway: video only (basic image)
- Canva: design + weak AI
- Pika: video only
- Adobe: requires multiple apps + subscriptions Craftit opportunity: First all-in-one AI-first creative platform. Canvas + auto-mode long video = unique.

Gap 2: The "Simplicity + Quality" Gap High-quality tools are complex. Simple tools produce mediocre quality.

- Runway/Luma: excellent quality, complex UX
- Pika/Haiper: simple UX, lower quality
- Canva: dead simple, weak AI Craftit opportunity: Runway-quality output with Pika-level simplicity. "Even a kid can do it" at high quality = unoccupied.

Gap 3: The "Auto Long Video" Gap Nobody has fully automated long-form video generation with editorial control.

- Runway: manual editing only
- Higgsfield Cinema Studio: manual 3D scene setup
- PixVerse: multi-shot but limited control
- OpenArt: One-Click Story is image-based, not video Craftit opportunity: Auto-mode (idea → script → scenes → camera → storyboard → final video) = first-mover advantage.

Gap 4: The "Design Canvas + AI Generation" Gap Canva has canvas but weak AI. AI tools have strong gen but no canvas.

- Canva: millions of templates, mediocre AI
- Midjourney/FLUX: great images, zero design tools
- Nobody combines powerful AI generation with design canvas Craftit opportunity: Canvas (image → posters, cards, invitations) fills exactly this gap.

Gap 5: The "Community + Creation" Gap Community features are either separate (Discord) or shallow.

- Midjourney: community is Discord, separate from creation
- Higgsfield: best in-app community but complex tool
- Canva/Adobe: no community
- Most: community is afterthought Craftit opportunity: Tightly integrated community. Creation → sharing → inspiration → remixing flywheel.

Gap 6: The "Mobile-First AI Creation" Gap Most AI tools are desktop-first or web-only.

- Runway: no mobile app
- Midjourney: no mobile app
- Luma: iOS but limited
- Canva: best mobile but weak AI Craftit opportunity: Killer mobile experience. Kids and casual users are mobile-first.

SECTION 5: USER PAIN POINTS FROM REDDIT (ranked list, red theme)

Create 15 pain points as ranked items:

#1 — Pricing & Credit Systems Feel Predatory (Intensity: 10/10) Credits burn fast, especially for video where failed/bad outputs still consume credits. Users feel punished for experimentation. Quote: "Every single feature is terrible. NEVER turns out well; waste of credits. Warped features, messed up hands." — r/RunwayML Quote: "Charging credits during beta is unacceptable." — r/RunwayML Quote: "Adobe doubled my price — so I finally cancelled. For what? Crappy AI features no one is asking for?" — r/VideoEditing (293 upvotes)

#2 — Poor Prompt Adherence & Unpredictable Output (Intensity: 9/10) Outputs described as "random generations" that don't follow instructions. Dozens of regenerations needed. Quote: "RANDOM generations, terrible prompt adherence, constantly changing face..." — r/RunwayML Quote: "Utterly random generations which make no sense at all." — r/RunwayML

#3 — Steep Learning Curve for Non-Technical Users (Intensity: 9/10) StableDiffusion/ComfyUI require Python/node-based workflows. Even Midjourney requires prompt engineering expertise. Quote: "It's mostly just telling the model directly what kind of skin you want — I mostly prompt for skin pores, detailed skin texture..." — r/midjourney user showing complex techniques beginners wouldn't know

#4 — Anatomical Artifacts — Hands, Fingers, Limbs (Intensity: 8/10) Persistent across all tools. Hands/fingers mangled. Bodies warp in video. Quote: "The hands and legs are utterly insanely mangled 99% of the gens; like it's SOOOOOOOO bad. COME ON Runway, are you guys running SD 1.5???" — r/RunwayML

#5 — AI Video Quality Still Far Below Expectations (Intensity: 8/10) Warped features, inconsistent characters, unrealistic motion, low resolution. Quote: "Video? Every single feature is terrible." — r/RunwayML Runway described as having "lost its lead" to Kling, Minimax, Sora.

#6 — Content Moderation Filters Too Aggressive (Intensity: 8/10) Legitimate creative projects blocked. Reasonable prompts rejected. Quote: "I keep getting blocked by the moderation system, claiming my prompt doesn't meet their standards — even when there's absolutely nothing inappropriate. At the price we're paying, why are they restricting us on BASIC creative needs?" — r/midjourney (true crime project)

#7 — Multi-Tool Workflow Fragmentation (Intensity: 7/10) 3-5+ tools per project: one for image gen, another for upscaling, video, editing, audio. Quote: "I prefer a flexible workflow where platforms combine several models — I don't like too many browser tabs opened." — r/generativeAI user who tested 15 tools Common workflow: MJ → Topaz → Runway → Premiere → ElevenLabs

#8 — Subscription Fatigue (Intensity: 7/10) $10-50/each across multiple platforms. Costs stack fast. One r/generativeAI comparison lists 15 tools priced $10-119/mo. Serious creators subscribe to 3-5 simultaneously.

#9 — Character/Face Consistency Across Generations (Intensity: 7/10) Can't maintain same character across images or video frames. Critical for storytelling, marketing, branding. Quote: "Constantly changing face" — common across r/RunwayML, r/midjourney Runway's likeness feature noted as "the ONLY product that has been decent" but still limited.

#10 — Lack of Fine Control Over Output (Intensity: 7/10) Users want specific camera angles, poses, lighting but lack precise controls. Runway's "multi-motion brush" and "fine-grain control" noted but users say these don't work well in practice.

#11 — Speed / Generation Time (Intensity: 6/10) Video gen takes minutes per attempt. When most fail, iteration cycle painfully slow. Fast rendering listed as competitive advantage for PixVerse. Users choose tools partly based on speed.

#12 — AI Content Labeling & Transparency Issues (Intensity: 6/10) AI elements mixed into stock libraries without labels. Quote: "Canva has been putting AI elements into their Element section, and they DO NOT show anything that warns if it's AI. I've been losing money because of this." — r/canva (421 upvotes, 0.98 ratio)

#13 — Resolution & Quality Limitations (Intensity: 6/10) Many tools output 480p-720p default. Upscaling adds another step/tool. Professional 4K expensive or unavailable. Budget tools like Pika output 480p on free tiers. 1080p premium, 4K rare.

#14 — Platform Lock-in & Walled Gardens (Intensity: 5/10) MJ originally Discord-only. Tools don't interoperate. Export limitations. Quote: "If the company you switch to accepts Visa/Mastercard, they will be forced to censor at some point." — r/StableDiffusion (2,269 upvotes) re: payment processor censorship

#15 — Lack of Audio Integration in Video Tools (Intensity: 5/10) AI video tools generate silent clips. Music, SFX, voiceover requires separate tools + manual sync. "Audio sync" differentiating feature for only a few tools (Veo 3.1). Most produce silent output.

SECTION 6: TOP 10 UNMET NEEDS (green theme)

#1 — All-in-One Creator Studio Single platform: image → video → audio → editing without switching tools. Quote: "We're working on a creator studio that bundles script generation, image/character creation, video generation, and audio tools into one workflow." — r/artificial (a startup attempting this, showing demand)

#2 — Consistent Characters Across Outputs Persistent character system maintaining identity across images, video clips, scenes. #1 feature request across r/midjourney, r/RunwayML, r/generativeAI.

#3 — Natural Language Control Without Prompt Engineering Plain language → result without specialized syntax, negative prompts, parameter tricks. Success of MJ prompt helper and PromptShot shows users need help translating intent.

#4 — Pay-Per-Success, Not Pay-Per-Attempt Pricing that doesn't charge for failed/bad generations. Pay for results, not wasted compute. Credit anxiety = major churn driver. Users afraid to experiment = worse outcomes = more churn.

#5 — One-Click Image-to-Video Pipeline Take generated/uploaded image → animate into coherent video with minimal effort. Multiple tools offer this but none consistently reliable.

#6 — Template-Based Creation for Non-Creators Pre-built templates: product ads, social posts, YouTube thumbnails, marketing videos. InVideo, HeyGen, Synthesia successful specifically because of templates. Canva users love templates.

#7 — Real-Time Preview / Live Generation Feedback See progress, course-correct before full render. Stop and redirect mid-generation. Currently: submit → wait → disappointed → submit again. No feedback loop.

#8 — Collaborative & Social Features for AI Creators Community specifically for AI-generated content. Share, feedback, learn. Quote: "We created a space where every single viewer is there specifically because they love AI-generated content. No algorithm fighting." — r/artificial AI content stigmatized on mainstream platforms. Creators want dedicated community.

#9 — Video Editing Integrated with Generation Generate clip → trim, combine, text, transitions, effects in same tool. Current: export raw AI video → import into Premiere/CapCut/DaVinci → edit. Major friction.

#10 — Automatic Translation & Localization Generate once → automatically adapt to multiple languages/cultural contexts. HeyGen's "auto video translation" is standout feature. Strong demand for global creation.

SECTION 7: SYNTHETIC USER PERSONAS (purple theme, 15 persona cards)

⚠️ Label this section: "SYNTHETIC RESEARCH — simulated personas, not real user data"

SEGMENT 1: Complete Beginners (3 personas)

Persona 1: Linda Chen, 42 — "Birthday Party Mom"

- Bio: Stay-at-home mom, 2 kids (6 and 9). Uses phone for WhatsApp, Facebook. Never used a design tool. Wants birthday invitations.
- First Encounter: Facebook mom-group post. Worried it's "like Photoshop."
- Onboarding: Expects sign up → tell it what she wants → get result. Confused by "prompt," "model," "resolution," "aspect ratio." Wants templates: "Birthday Invitation," "Thank You Card."
- First Creation: "Pink princess birthday invitation for my daughter Emma, turning 7." Hard if: needs to figure out aspect ratio, text doesn't render correctly (names/dates).
- Repeat Usage: Holiday cards, school event flyers. Leaves if: every creation from scratch, too AI-looking.
- Pain Points: Text not rendering correctly, no edit path, overwhelmed by options, worried about cost.
- Delight: First "professional" image — "I made THIS?!" Direct WhatsApp sharing. Canvas for text overlays.
- WTP: $3-5/mo. Free tier essential. Would pay one-time for "party pack."
- Magic Moment: "I described what I wanted and got something PROFESSIONAL."

Persona 2: Rajesh Patel, 55 — "Local Shop Owner"

- Bio: Small electronics repair shop. Basic smartphone. Posts occasionally on FB/IG with stock photos.
- First Encounter: Nephew says "Uncle, try this app, it makes ads for you."
- Onboarding: Expects to describe shop → get content. Needs Hindi language support. Wants "Social Media Post" templates.
- First Creation: "Professional ad for mobile phone repair shop, fast service." Hard: making it look like HIS shop, adding phone number/address.
- Repeat Usage: Weekly social posts, festival announcements, WhatsApp status. Leaves if: can't customize branding, takes too long.
- Pain Points: Can't upload logo, text in local language renders poorly, results generic, no batch creation.
- Delight: Post gets more likes. Video ad looks "like a real commercial." Festival greeting with shop branding.
- WTP: $2-4/mo. Very price-sensitive. Free tier essential.
- Magic Moment: First post with noticeably more engagement.

Persona 3: Aisha, 16 — "Fun Content Teen"

- Bio: High schooler. Lives on TikTok/IG. Loves memes, aesthetic edits, fan art. Phone only.
- First Encounter: Viral TikTok of AI art. "Omg I need this." Downloads immediately.
- Onboarding: Zero patience for tutorials. Wants creation within 30 seconds. Skip anything longer than 2 taps.
- First Creation: "My cat as a samurai in space." Loses interest if >15-20 seconds generation.
- Repeat Usage: Memes, fan edits, aesthetic profile pics, TikTok content. Leaves if: friends move on, free tier too limited, watermarks.
- Pain Points: Watermarks (social stigma), wait times, can't make video memes, moderation blocking harmless prompts.
- Delight: Making something so cool friends ask "HOW?" Video from own photos. Community likes/follows.
- WTP: $0. Might convince parents for $5/mo. Would watch ads for credits.
- Magic Moment: "Everyone wants to know how I made that."

SEGMENT 2: Casual Creators (3 personas)

Persona 4: Sophie Müller, 29 — "Canva Regular"

- Bio: Marketing coordinator. Uses Canva daily for social, presentations. Comfortable with templates, not a designer.
- First Encounter: "Best AI Design Tools 2026" blog. "Can this replace some of my Canva workflow?"
- Onboarding: Expects professional onboarding, workspace, brand kit upload. Compares everything to Canva UX.
- First Creation: Social media post — product announcement. Hard: matching brand guidelines, text overlays, exact social dimensions.
- Repeat Usage: Unique hero images, campaign visuals. Leaves if: can't integrate into workflow, no brand consistency.
- Pain Points: No brand kit, can't output specific social sizes, no text editing layer, inconsistent style.
- Delight: 10 campaign variations in 5 minutes vs hours. Canvas bridging AI→finished design. "Style lock."
- WTP: $10-15/mo. Would pitch manager for team plan at $20-30/mo.
- Magic Moment: "5 minutes instead of 2 hours in Canva."

Persona 5: Marcus Johnson, 34 — "CapCut Hobbyist"

- Bio: Hobbyist travel vlogger. Uses CapCut, Premiere Rush. Wants AI for B-roll and transitions.
- First Encounter: YouTube video "I used AI to generate B-roll." "Quality needs to be good enough for real footage."
- Onboarding: Expects video-first. Confused if app feels image-centric.
- First Creation: 8-second aerial tropical beach B-roll. Hard: matching look/feel of footage, camera control, 4K.
- Repeat Usage: B-roll gaps, establishing shots, transitions. Leaves if: noticeably "AI-looking," no camera control.
- Pain Points: Quality/resolution, no camera movement control, can't match color grade, 4-12s limit.
- Delight: AI B-roll seamlessly cutting into real footage. First/last frame for transitions.
- WTP: $15-20/mo. Values quality over quantity.
- Magic Moment: "AI B-roll that seamlessly cuts into real footage."

Persona 6: Priya Sharma, 26 — "Instagram Aesthetic Curator"

- Bio: HR by day. Runs curated IG aesthetic account (12K followers). Uses VSCO, Lightroom Mobile, Canva.
- First Encounter: Sees #CraftitAI tag on stunning surreal landscape. DMs creator.
- Onboarding: Expects style-focused. Wants to define aesthetic and maintain it.
- First Creation: Dreamy pastel surreal landscape in her signature aesthetic. Hard: matching her very specific Instagram aesthetic.
- Repeat Usage: Daily posts, story content, unique visuals. Leaves if: others use same styles, content looks generic.
- Pain Points: Style consistency across sessions, others generating identical content, no "style recipe," community exposing her source.
- Delight: Custom style with pixel-perfect consistency. Full week of grid content in one session. Video versions for Reels.
- WTP: $8-12/mo. Would pay premium for exclusive styles.
- Magic Moment: "Full week of cohesive grid content in one session."

SEGMENT 3: Content Creators / Influencers (3 personas)

Persona 7: Jake Torres, 27 — "YouTube Thumbnail Machine"

- Bio: YouTube 180K subs, 3x/week. Spends too much time on thumbnails. Uses PS + stock.
- First Encounter: Creator podcast. "If this cuts thumbnail time from 30 to 5 min, I'm in."
- Onboarding: Wants power features upfront. Batch creation, templates, instant export. Needs ROI in first session.
- First Creation: Dramatic thumbnail. Hard: incorporating his actual face, consistent character, text integration.
- Repeat Usage: Every thumbnail, A/B testing. Leaves if: can't use his face consistently, quality doesn't match top YT thumbnails.
- Pain Points: Can't use own likeness reliably, no text overlay tools, no A/B testing, style drift hurts brand.
- Delight: 5 thumbnail variations in 2 minutes. Canvas for text. Higher CTR.
- WTP: $20-30/mo without hesitation. Business expense.
- Magic Moment: "5 thumbnail variations in 2 minutes."

Persona 8: Mia Zhang, 24 — "TikTok Content Factory"

- Bio: Full-time TikTok 500K followers. 2-3 posts daily across TikTok/Reels/Shorts. Small team (1 editor).
- First Encounter: Creator DMs her. "Can this keep up with my volume? I need FAST."
- Onboarding: Speed above all. Bulk gen, API, preset reuse. Confused if one-at-a-time design.
- First Creation: Series of 6-second clips for trending audio. Hard: batch creation, style consistency, matching trends fast.
- Repeat Usage: Daily production. Leaves if: too slow, competitors offer API/bulk, quality plateaus.
- Pain Points: No batch/bulk, can't save prompt templates, speed not fast enough for trends, no posting integration.
- Delight: Auto-mode idea→finished video. Full day's content in 20 minutes. Reference image for series consistency.
- WTP: $30-50/mo for unlimited/high-volume. Would pay for API.
- Magic Moment: "Full day's content in 20 minutes."

Persona 9: David Okafor, 31 — "Multi-Platform Brand Builder"

- Bio: Personal brand — LinkedIn + IG + YouTube. Entrepreneurship content. Consistent visual branding needed.
- First Encounter: LinkedIn post about AI content creation.
- Onboarding: Expects brand kit upload, multi-format export (LinkedIn banner, IG square, YT thumbnail from same design).
- First Creation: Motivational quote image for LinkedIn + IG. Hard: specific brand colors/fonts, resizing.
- Repeat Usage: Daily cross-platform content. Leaves if: inconsistent branding, no multi-format.
- Pain Points: No multi-format export, brand re-prompting, no visual library, no scheduling.
- Delight: One prompt → 3 platforms correct formats. Locked brand style. Auto-mode video from blog post.
- WTP: $25-40/mo. Brand consistency is key value driver.
- Magic Moment: "One prompt → content for 3 platforms."

SEGMENT 4: Small Business Owners (2 personas)

Persona 10: Emma Larsson, 38 — "E-commerce Product Seller"

- Bio: Handmade jewelry on Etsy/IG. Does own photography (mediocre). Needs better product visuals + video ads.
- First Encounter: Etsy seller forum: "This AI tool made my product photos look amazing."
- Onboarding: Expects upload product photo → get professional variations. Needs "product photography" workflow.
- First Creation: Upload necklace photo → wants it on model in luxury setting. CRITICAL: product must look exactly right, AI must not alter jewelry design.
- Repeat Usage: Seasonal campaigns, product launches, video ads. Leaves if: product accuracy unreliable.
- Pain Points: AI changes product details, can't generate "model wearing product," no e-commerce templates, needs lifestyle context with real products.
- Delight: Reference image preserving product in changed scene. Video ad looking professional. Canvas for sale graphics.
- WTP: $10-20/mo. Direct ROI: better photos → more sales.
- Magic Moment: "Video ad that looks like a professional commercial."

Persona 11: Carlos Mendez, 45 — "Restaurant Owner"

- Bio: Family Mexican restaurant. Manages own social media. Phone food photos. Zero design skills.
- First Encounter: Google "easy way to make restaurant social media posts."
- Onboarding: Expects "I'm a restaurant" → templates and workflows. Confused by any creative decision-making.
- First Creation: "Taco Tuesday special, $5 tacos." Problem: AI tacos won't look like HIS tacos. Needs reference images.
- Repeat Usage: Weekly specials, events, seasonal menus, Google Business. Leaves if: >5 min per post, obviously AI-generated.
- Pain Points: AI food doesn't match real dishes, no calendar/planning, uncanny AI food, needs Spanish support.
- Delight: Reference image → professional food photography. 12-second video ad. One-tap post to IG.
- WTP: $5-10/mo. Must be dead simple. Would pay more for "restaurant marketing" package.
- Magic Moment: "Reference image turns his tacos into professional food photography."

SEGMENT 5: Professional Creators (2 personas)

Persona 12: Alex Kim, 30 — "Freelance Graphic Designer"

- Bio: 6 years freelance. Figma, Illustrator, PS daily. Evaluating AI to speed up ideation. Skeptical.
- First Encounter: Design Twitter comparing AI tools. "Let's see if this is useful or just a toy."
- Onboarding: Expects pro controls. Understand model, adjust parameters, control composition. Frustrated by oversimplified UIs.
- First Creation: Concept art for client — specific composition, color palette, mood. Easy: rapid mood boards. Hard: precise compositional control, specific palettes, print resolution.
- Repeat Usage: Ideation, concept variations, generating assets to refine in Figma/PS. Leaves if: no useful exports (PSD, SVG, hi-res PNG), feels like toy.
- Pain Points: Insufficient composition control, no layer separation, low resolution for print, no API, feels "dumbed down."
- Delight: 20 concept directions in 10 minutes for client review. Canvas bridging AI→design. Image-to-video for motion concepts.
- WTP: $20-30/mo for pro tier. Would pay for hi-res, advanced controls, API, commercial licensing clarity.
- Magic Moment: "20 concept directions in 10 minutes for client review."

Persona 13: Nina Volkov, 35 — "Video Production Lead"

- Bio: Leads 5-person video production company. Commercials, corporate, social for clients. Exploring AI for pre-viz and storyboarding.
- First Encounter: Industry conference demo. "Quality isn't broadcast-ready yet, but storyboarding could be valuable NOW."
- Onboarding: Expects project workflow, team collaboration, high quality.
- First Creation: Storyboard for commercial — 8 scenes, specific camera angles, consistent characters. Hard: consistency across scenes, camera angles, visual continuity.
- Repeat Usage: Pre-viz, storyboarding, pitch visuals, B-roll. Leaves if: no consistency, no team features.
- Pain Points: Character/scene consistency critical and lacking, no storyboard view, can't specify camera angles, no team features, manual timeline tedious.
- Delight: Auto-mode script→storyboard→video = game-changing. First/last frame for continuity. Storyboard in 30 min vs 2 days.
- WTP: $50-100/mo team. Justifies as production cost savings.
- Magic Moment: "Storyboard deck in 30 minutes vs 2 days."

SEGMENT 6: Kids / Students (2 personas)

Persona 14: Tommy, 11 — "School Project Kid"

- Bio: 6th grader. Needs visuals for school presentations. Uses Google Slides.
- First Encounter: Classmate shows at school. "This is SO COOL, I want a dinosaur fighting a robot."
- Onboarding: Type something → get something cool. Zero patience. Kid-friendly interface. Parents need: content safety, privacy.
- First Creation: "Roman gladiator in colosseum for history project." Frustrated if: content filter blocks "gladiator."
- Repeat Usage: Every school project, showing friends, exploring creativity. Leaves if: free credits run out.
- Pain Points: Aggressive moderation on school topics, parent concerns about academic integrity, limited free tier, can't use in Google Slides easily.
- Delight: First "whoa" showing friends/teachers. Video for presentation. Community for school project art.
- WTP: Parents $3-5/mo for "education" tier. School licensing avenue.
- Magic Moment: "Everyone in class wanted to know how I made my presentation look so cool."

Persona 15: Zara, 19 — "Art School Student"

- Bio: First-year art student. Procreate, some PS. Interested in AI but conflicted about ethics.

- First Encounter: Professor assigns "create with AI tools and critically analyze."
- Onboarding: Wants creative control. Interested in process, not just output.
- First Creation: Surrealist self-portrait — uploads reference sketch. Hard: using her sketches as starting points, maintaining artistic intent.
- Repeat Usage: Ideation, style exploration, reference material, assignments. Leaves if: tool does ALL creative work, ethical concerns increase.
- Pain Points: No sketch-to-image workflow, ethical concerns about training data, no control over how much AI changes vs preserves, watermark/attribution unclear.
- Delight: Reference image with her sketch as starting point. AI as collaborator. Video art blending AI + traditional.
- WTP: $5-8/mo with student discount. Values ethical transparency.
- Magic Moment: "AI as collaborator, not replacement — keeping her style while expanding."

Cross-Segment Synthesis (10 key insights as connected sticky notes):

1. **The 30-Second Rule** — Every segment expects wow within 30-60 seconds. Any friction before first creation is fatal.
2. **Text on Images is Universal** — From birthday invitations to YouTube thumbnails. Canvas is essential, not nice-to-have.
3. **Consistency = #1 Professional Barrier** — Style/character/brand consistency cited by segments 3-5 as top pain point.
4. **Reference Image is the Killer Feature** — Using own images as starting points separates "cool AI toy" from "useful for MY life."
5. **"Is This My Work?" Identity Question** — Users want to feel ownership. Tool must amplify THEM, not replace them.
6. **Price Sensitivity Spans 50x** — $0 (teens) to $100/mo (production teams). Single tier won't work.
7. **Templates > Blank Canvas** for 80% of users. "What do you want to make?" → category → template → customize.
8. **Video is Differentiator, Images are Entry** — Users discover for images, stay for video. Auto-mode long video is moonshot.
9. **Community is Double-Edged** — Teens love it (validation). Curators fear it (exposing secrets). Pros ignore it. Must be opt-in.
10. **Speed = Respect** — Every persona mentions speed. Teen = impatience. Business = "no time." Creator = "need volume."

Segment Needs Matrix:

Need	Beginners	Casual	Creators	SMB	Pro	Kids
Templates/guided flows	★★★	★★	★	★★★	☆	★★
Text on images	★★★	★★★	★★★	★★★	★★	★
Brand consistency	☆	★★	★★★	★★	★★★	☆
Reference image	★	★★	★★	★★★	★★★	★
Batch/bulk creation	☆	★	★★★	★★	★★	☆
Video generation	★	★★	★★★	★★	★★★	★★
Multi-format export	☆	★★	★★★	★★	★★★	☆
Simple/fast UX	★★★	★★	★★	★★★	★	★★★
Advanced controls	☆	★	★★	☆	★★★	☆
Community/social	★	★★	★	☆	☆	★★★
Auto-mode video	★	★	★★★	★★	★★★	★
Content safety	★★	★	★	★	★	★★★

★★★ = Critical | ★★ = Important | ★ = Nice to have | ☆ = Low priority

Critical UX Moments (Make-or-Break):

1. **First Generation Result** — If first image/video isn't impressive, users leave. Optimize default quality, pre-fill good prompts.
2. **"How Do I Edit This?"** — After generation, user wants to change ONE thing. No edit path = stuck. Canvas addresses this.
3. **"It Doesn't Look Like MY Brand/Product"** — Professional users realize output is generic. Reference images + style saving + brand kits needed.
4. **Sign-Up Wall Timing** — Requiring sign-up BEFORE first generation = conversion killer. Let them generate one free, then gate.
5. **Credit Exhaustion** — When free credits run out. Pricing, value communication, upgrade UX either convert or lose forever.
6. **Share Flow** — Every extra tap between "done" and "shared" is friction. One-tap share essential.
7. **Second Session Return** — Day 2-3. Does user have reason to come back? Saved projects, templates, "try this next" suggestions.

Feature Priority (User Perspective):

Tier 1: Must-Have (blocks adoption)

1. Canvas / text overlay on AI images
2. Reference image generation (use my photos as input)
3. Simple guided creation flows (templates by use case)
4. Fast generation (< 30 seconds for images)
5. Easy sharing (one-tap to social/messaging)

Tier 2: Growth Drivers (converts trial → paid)

6. Style saving & consistency across sessions
7. Brand kit (colors, fonts, logo)
8. Multi-format export (social media sizes)
9. Video generation quality improvements
10. Auto-mode long video (idea → final video)

Tier 3: Retention & Expansion

11. Batch/bulk generation
12. Team collaboration
13. API access
14. Advanced controls (composition, camera, parameters)
15. Scheduling / publishing integration

Magic Moments Per Segment:

Segment	Magic Moment
Complete Beginners	"I described what I wanted and got something PROFESSIONAL. I made this!"
Casual Creators	"I just made in 5 minutes what used to take me 2 hours in Canva."
Content Creators	"I generated a full week of content in one sitting."

Segment	Magic Moment
Small Business	"A customer said my ad looked like it was made by a professional agency."
Professional	"I showed 10 concept directions in the first meeting — they were blown away."
Kids/Students	"Everyone in class wanted to know how I made my presentation look so cool."

SECTION 8: STRATEGIC OPPORTUNITIES (ranked, with evidence)

🥇 **#1 — Own the "Full Pipeline" Position** Be the first all-in-one AI creative platform: image → canvas → short video → long video. Evidence: Competitor analysis = zero tools with full pipeline. Reddit #1 unmet need = all-in-one studio. All synthetic personas want fewer tools.

🥇 **#2 — Auto-Mode Long Video = Category-Defining** Idea → script → scenes → camera → storyboard → final video with edit-at-every-step. Evidence: PixVerse/OpenArt closest but neither has full vision. Nina (pro persona) called this "game-changing." Creators pay $30-50/mo for this alone.

🥇 **#3 — "Simplicity + Quality" = Unoccupied** Runway-quality with Pika-simplicity. "Even a kid can do it" at high quality. Evidence: Reddit #3 pain = learning curve. Synthetic: 30-second rule. Canva proved model at scale (260M MAU).

🥈 **#4 — Canvas + AI = Canva Killer** Canva has canvas with weak AI. AI tools have strong gen with no canvas. Combine both. Evidence: Text on images = universal need across all segments. Canvas = "essential, not nice-to-have."

🥈 **#5 — Mobile-First = Blue Ocean** 3B+ smartphone users. SEA, India, LatAm are mobile-only markets. Evidence: Runway, MJ = no mobile app. Asia-Pacific growing 27.6% CAGR. Target audience (kids, beginners) = mobile-first.

🥈 **#6 — Fair, Transparent Pricing** Flat monthly. Don't charge for failures. Evidence: Reddit #1 pain. Credit anxiety = major churn. Users afraid to experiment = worse outcomes = more churn.

🥉 **#7 — Kids & Family (Long-term)** 50-100M kids. Zero competition. COPPA = barrier others won't cross. Evidence: Market research: zero tools. Canva Education 60M+ students. Parents pay $3-15/mo.

🥉 **#8 — Community as Growth Engine** Contests, remixing, galleries. Creation → sharing → inspiration flywheel. Evidence: MJ built on Discord (19M). Higgsfield runs $500K contests. AI creators want dedicated community.

SECTION 9: PIVOT RECOMMENDATION (blue theme)

Current Position:

Craftit AI positioned as "AI content generation tool" — one of many in crowded, fragmenting market heading toward consolidation.

Recommended Pivot:

FROM: AI content generation tool TO: The AI-first creative platform for everyone

Reframing:

Old Framing	New Framing
Image Generation	"What do you want to make?" → templates by use case
Video Generation	Seamless part of creation flow, not separate feature
Canvas (upcoming)	THE core product experience — where generation meets design
Auto Long Video (upcoming)	Category-defining — idea to finished video
Community	Growth engine — creation → sharing → inspiration → creation

Positioning Statement:

"Craftit AI is the AI-first creative platform that makes it dead simple for anyone — from kids to professionals — to create stunning images, videos, and designs. Unlike Canva (great design tools, weak AI), Midjourney (great images, no design tools), or Runway (great video, complex UX), Craftit AI combines powerful AI generation with intuitive design tools in a single, simple platform."

Key Differentiators (in defensibility order):

1. Full pipeline — image → canvas → short video → long video (NO competitor does this)
2. Auto-mode long video — idea → script → storyboard → final video (first mover)
3. Simplicity at quality — Runway output, Pika simplicity
4. Canvas + AI — Canva's design power with MJ-level generation

Target Segments (in order):

1. NOW: Casual creators + content creators ($8-30/mo, most frustrated, most vocal)
2. NEXT: Small businesses + beginners (template-driven, massive volume)
3. LATER: Kids/Education + mobile-first markets (blue ocean, zero competition)

SECTION 10: ROADMAP (timeline view with 3 phases)

Phase 1: NOW (0-4 weeks) — Quick Wins

Action	Impact	Effort
Template-first entry points — "Birthday Card," "Product Ad," "Social Post" not "Image Generation"	Reduces beginner drop-off	Low — UI reorg
AI-assisted prompting — simple words in, enhanced prompt behind scenes	Addresses learning curve pain #3	Medium
Transparent pricing page — flat monthly, no credit anxiety, generous free tier	Addresses predatory pricing pain #1	Low — pricing decision
One-tap social sharing — Done → WhatsApp/IG/TikTok	Every share = free marketing	Medium
Rename/reframe — "AI Creative Platform" not "AI Image Generator"	Signals full-pipeline ambition	Low — branding

Phase 2: NEXT (1-3 months) — Ship the Moat

Feature	Why It's a Moat	Priority
Content Creation Canvas	Combines AI gen + design. No competitor has both at quality. Makes everything more valuable.	🔴 CRITICAL
Auto-mode long video	Category-defining. Idea → final video. First-mover.	🔴 CRITICAL
Style/brand saving	Converts casual → paying. Critical for creators/business.	🟡 HIGH
Character persistence	Unlocks storytelling, branding, series content.	🟡 HIGH
Reference image improvements	Bridges "cool toy" to "useful for MY life."	🟡 HIGH

Phase 3: LATER (3-6 months) — Scale

Feature	Why	Priority
Mobile-first app	Blue ocean. Target audience is mobile-first.	🟡 HIGH
Community 2.0	Contests, remixing, galleries. Growth flywheel.	🟠 MEDIUM
Batch/bulk generation	Unlocks high-volume creators ($30-50/mo).	🟠 MEDIUM
Multi-format export	One design → LinkedIn, IG, YT sizes auto.	🟠 MEDIUM
Team/collaboration	Shared brand kits, multi-seat. Opens SMB/enterprise.	🟠 MEDIUM
Kids/Education tier	Safe, moderated. School partnerships.	🔵 STRATEGIC
API access	Power users, developer ecosystem.	🔵 STRATEGIC

SECTION 11: PRICING RECOMMENDATION

Tier	Price	Target	Includes
Free	$0	Everyone	Generous gens, no embarrassing watermarks
Starter	$7-9/mo	Beginners, students, casual	More gens, all basics, HD export
Pro	$15-19/mo	Creators, small business	Unlimited gen, brand kit, style saving, priority speed, auto-mode video
Team	$35-49/mo	Teams, agencies	Multi-seat, shared brand, collaboration, bulk gen

Principles:

- Undercut Runway ($76/mo unlimited) dramatically
- Match or beat Pika ($8-28/mo)
- More generous than Canva Pro ($13/mo) on AI features
- NEVER charge for failed generations

- Annual discount: 20-30%

SECTION 12: RISKS & MITIGATIONS

Risk	Likelihood	Mitigation
Canva improves AI quality	HIGH	Ship canvas + auto-video first. Be AI-native not AI-bolted-on. Speed matters.
Model quality becomes commodity	HIGH	Compete on workflow/UX, not model alone. Full pipeline is the moat.
Runway simplifies UX	MEDIUM	Move faster on simplicity. Runway's DNA = pro; Craftit's DNA = everyone.
AI regulation tightens	MEDIUM	Build content credentials early. Turn compliance into feature.
Cash runway constraints	HIGH	Focus on virality (sharing, community) for organic growth. Keep lean.

SECTION 13: METRICS

North Star: Creations shared per week

Supporting:

Category	Metric	Target
Activation	% users who create within 60 seconds	>60%
Engagement	Creations per user per week	>3
Retention	Week 1 retention	>40%
Conversion	Free → paid	>5%
Growth	Creations shared externally per week	+15% MoM
Revenue	ARPU	>$12/mo
Satisfaction	NPS	>50

SECTION 14: IMMEDIATE NEXT STEPS

1. Align on pivot direction — "AI-first creative platform for everyone"
2. Prioritize Canvas launch — #1 strategic feature
3. Redesign entry experience — template-first, use-case-first
4. Set pricing — flat monthly, generous free, no credit anxiety
5. Plan auto-mode long video architecture
6. Run 5-8 real user interviews to validate synthetic findings

SECTION 15: KEY COMPETITIVE LANDSCAPE (summary table)

Competitor	Users	Revenue	Strength	Weakness
Canva	260M MAU	$3.5B	UX, templates, scale	Not AI-native; complex for kids
Midjourney	20M registered	$500M	Quality, community	Desktop-heavy, prompt-dependent
Adobe Firefly	30M+ (CC)	Part of $20B+	Pro integration	Expensive, complex, enterprise
Leonardo AI	~10M+	~$50-100M est.	Fine-tuning, gaming art	Niche, complex
Runway	~5M+	~$50-100M est.	Video leader	Pro-focused, expensive
Pika	~5M+	~$30-50M est.	Simple video gen	Early stage, limited
Ideogram	Growing fast	Undisclosed	Text-in-image, quality	Narrow focus

SECTION 16: POTENTIAL PARTNERSHIPS

Partner Type	Examples	Value
Education platforms	Google Classroom, Khan Academy, Kahoot	Distribution to 100M+ students
Children's media	Disney, Nickelodeon, PBS Kids	Trusted brands + content IP
Social platforms	TikTok, Instagram, YouTube Shorts	In-app creation integration
E-commerce	Shopify, Etsy, Amazon Merch	AI-generated product designs
Print-on-demand	Printful, Redbubble, Zazzle	Monetization of creations
Device OEMs	Samsung, Xiaomi, Apple	Pre-installed app deals
Telecoms	Jio (India), Telkomsel (Indonesia)	Bundled subscriptions emerging markets
School districts	US K-12, international schools	Institutional sales

TIMING ADVANTAGE (prominent callout)

Ship Canvas + Auto-mode Long Video in 2026 and own the "complete AI creative suite for everyone" position BEFORE Canva catches up on AI quality or Runway simplifies their UX.

End of data. Please create this as a comprehensive, well-organized, visually rich FigJam board with all sections clearly separated, color-coded, and connected where relevant. Use sticky notes for individual items, sections/frames for grouping,

and connectors to show relationships between insights.

PROMPT END